

- Adversarial search (Ch 5 of R-N)
- Also see Chapter 12 of Edelkamp.
- Two-player games

# Two-player games overview

- Computer programs which play 2-player games
  - game-playing as search
  - with the complication of an opponent
- General principles of game-playing and search
  - evaluation functions
  - minimax principle
  - alpha-beta-pruning
  - heuristic techniques


# Two-player games overview

- Status of Game-Playing Systems
  - in chess, checkers, backgammon, Othello, etc, computers routinely defeat top human players (Checkers has been shown to be a draw when played by both optimally.)
- Applications?
  - think of “nature” as an opponent
  - economics, war-gaming, medical drug treatment

# Games of strategy

- Deterministic rules (or deterministic rules plus probabilistic rules – these are games that combine strategy and luck, e.g. bridge, backgammon, blackjack).
- Moves are alternately made by two players A and B.
- rules define how configurations change
- A subset F of configurations is identified as final.
- typically F is partitioned into three sets: T, A and B.
  - T is tie, A (B) is win for player A (B)
- Goal is to develop a strategy for one player to win. (computer plays for that player)

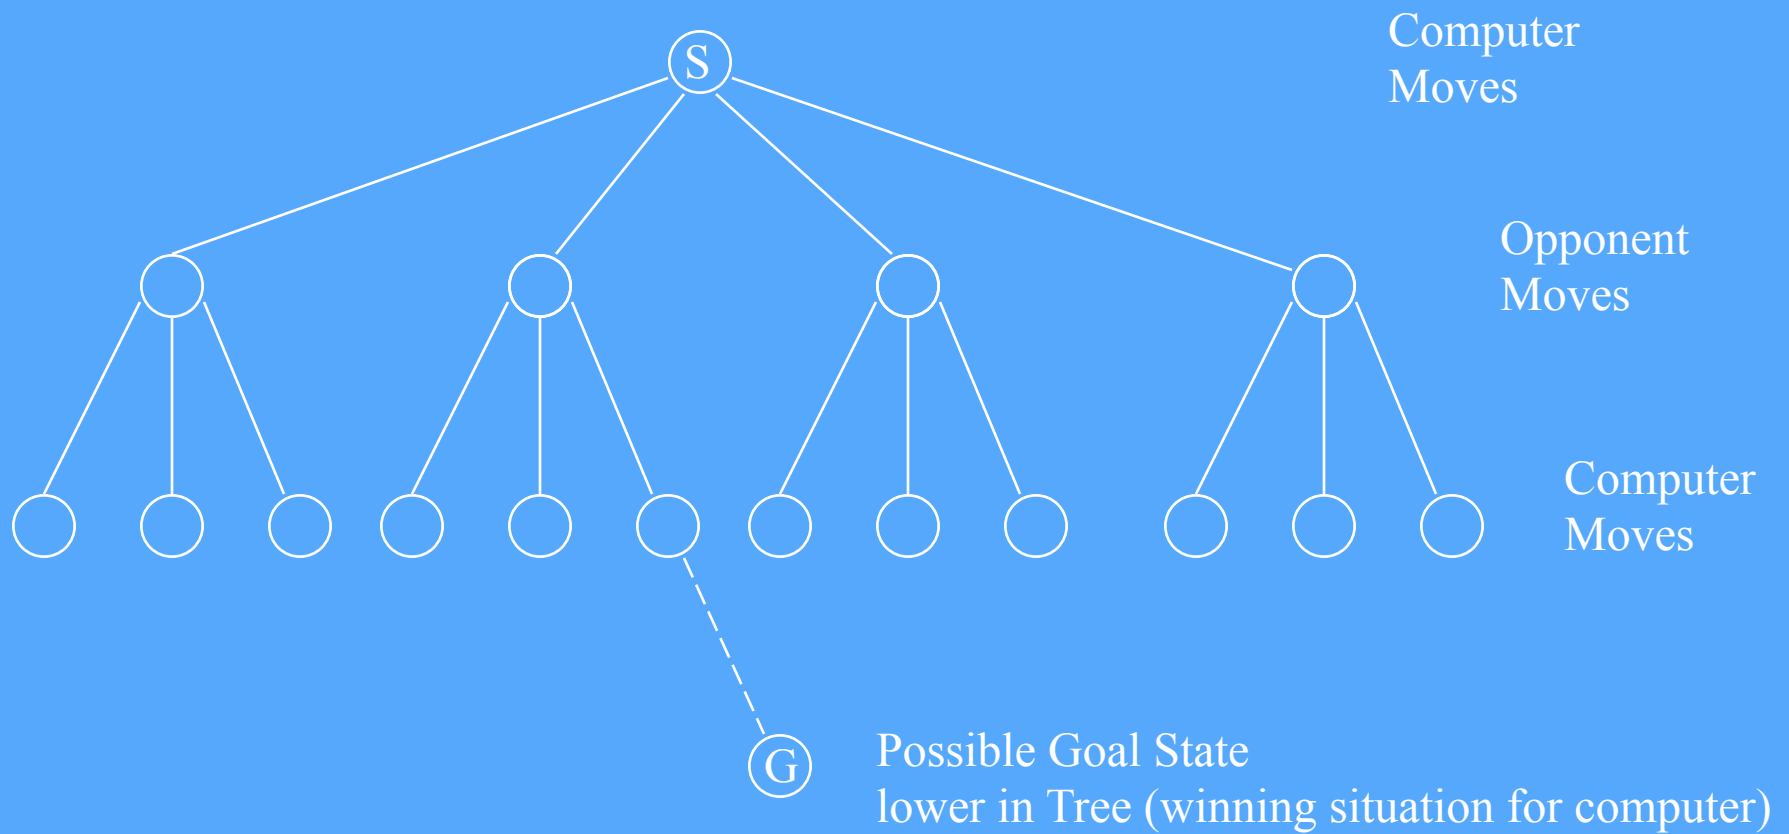
# Chess Rating Scale



# Two-Player Games with Complete Trees


- Let us assume that the computer has the first move. Then, the game can be described as a **series of decisions**, where the first decision is made by the computer, the second one by the human, the third one by the computer, and so on, until all coins are gone.
- The **computer** wants to make decisions that **guarantee its victory** (in this simple game).
- The underlying assumption is that the **human** always finds the **optimal move**.

# Game Tree Representation



- New aspect to search problem
  - there's an opponent we cannot control
  - how can we handle this?

# Game Trees



**Figure 5.1** A (partial) search tree for the game of Tic-Tac-Toe. The top node is the initial state, and MAX moves first, placing an X in some square. We show part of the search tree, giving alternating moves by MIN (O) and MAX, until we eventually reach terminal states, which can be assigned utilities according to the rules of the game.



# Min-max algorithm

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$ 
```

---

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return v
```


---

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

---

**Figure 5.3** An algorithm for calculating minimax decisions. It returns the action corresponding to the best possible move, that is, the move that leads to the outcome with the best utility, under the assumption that the opponent plays to minimize utility. The functions MAX-VALUE and MIN-VALUE go through the whole game tree, all the way to the leaves, to determine the backed-up value of a state. The notation  $\arg \max_{a \in S} f(a)$  computes the element *a* of set *S* that has the maximum value of *f(a)*.

# Game Trees



**Figure 5.2** A two-ply game tree as generated by the minimax algorithm. The  $\triangle$  nodes are moves by MAX and the  $\nabla$  nodes are moves by MIN. The terminal nodes show the utility value for MAX computed by the utility function (i.e., by the rules of the game), whereas the utilities of the other nodes are computed by the minimax algorithm from the utilities of their successors. MAX's best move is  $A_1$ , and MIN's best reply is  $A_{11}$ .

# Min-max Principle

- “Assume the worst”
  - say each configuration has an evaluation number
  - high numbers favor the player (the computer)
    - so we want to choose moves which maximize evaluation
  - low numbers favor the opponent
    - so they will choose moves which minimize evaluation
- Minimax Principle
  - you (the computer) assume that the opponent will choose the minimizing move next (after your move)
  - so you now choose the best move under this assumption
    - i.e., the maximum (highest-value) option considering both your move and the opponent's optimal move.

# An optimal procedure: The Min-Max method

- Designed to find the optimal strategy for Max and find best move:
  - 1. Generate the whole game tree to leaves
  - 2. Apply utility (payoff) function to leaves
  - 3. Back-up values from leaves toward the root:
    - a Max node computes the max of its child values
    - a Min node computes the Min of its child values
  - 4. When value reaches the root: choose max value and the corresponding move.
- However: It is impossible to develop the whole search tree, instead develop part of the tree and evaluate promise of leaves using a static evaluation function.

# More general Minmax search algorithm

## Procedure Minimax

**Input:** Position  $u$

**Output:** Value at root


```
if (leaf( $u$ )) return  $Eval(u)$                                 ;; No successor, static evaluation
if (max-node( $u$ ))  $val \leftarrow -\infty$                         ;; Initialize return value for MAX node
else  $val \leftarrow +\infty$                                     ;; Initialize return value for MIN node
for each  $v \in Succ(u)$                                        ;; Traverse successor list
    if (max-node( $u$ ))  $val \leftarrow \max\{val, Minimax(v)\}$     ;; Recursive call at MAX node
    else  $val \leftarrow \min\{val, Minimax(v)\}$                 ;; Recursive call at MIN node
return  $res$                                                   ;; Return final evaluation
```

## Algorithm 12.2

---

Minimax game tree search.

What is the minimax value of the root node?



# Complexity of Game Playing

- Imagine we could predict the opponent's moves given each computer move
- How complex would search be in this case?
  - worst case, it will be  $O(b^d)$
  - Chess:
    - $b \sim 35$  (average branching factor)
    - $d \sim 100$  (depth of game tree for typical game)
    - $b^d \sim 35^{100} \sim 10^{154}$  nodes!!
  - Tic-Tac-Toe
    - ~5 legal moves, total of 9 moves
    - Number of states ~ 26000 (counting all rotations, reflections ...)
- well-known games can produce enormous search trees

# Static (Heuristic) Evaluation Functions

- An Evaluation Function:
  - estimates how good the current board configuration is for a player.
  - Typically, one figures how good it is for the player, and how good it is for the opponent, and subtracts the opponents score from the players
  - Othello: Number of white pieces - Number of black pieces
  - Chess: Value of all white pieces - Value of all black pieces
- Typical values from -infinity (loss) to +infinity (win) or  $[-1, +1]$ .
- If the board evaluation is  $X$  for a player, it's  $-X$  for the opponent





# Two-Player Games Example: Tic-tac-Toe


$$e(p) = 8 - 8 = 0$$

X		
O	X	

$$e(p) = 6 - 2 = 4$$

O	O	X
X	O	
X		

$$e(p) = 2 - 2 = 0$$

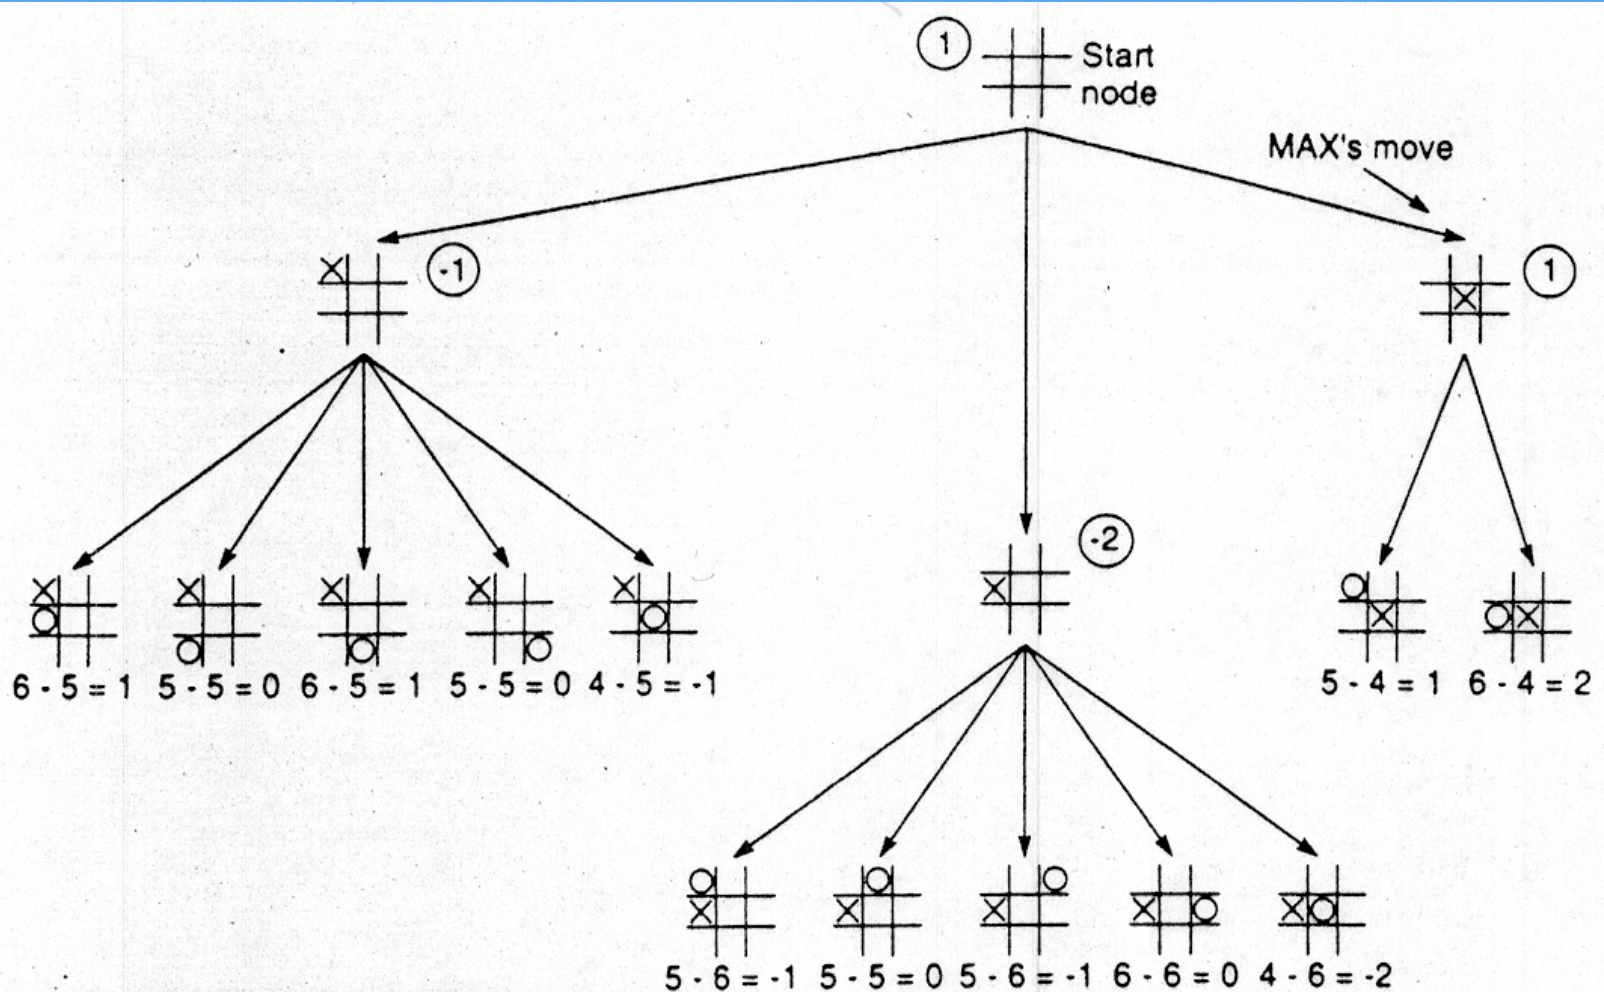
O	O	X
	X	
X		

$$e(p) = \infty$$

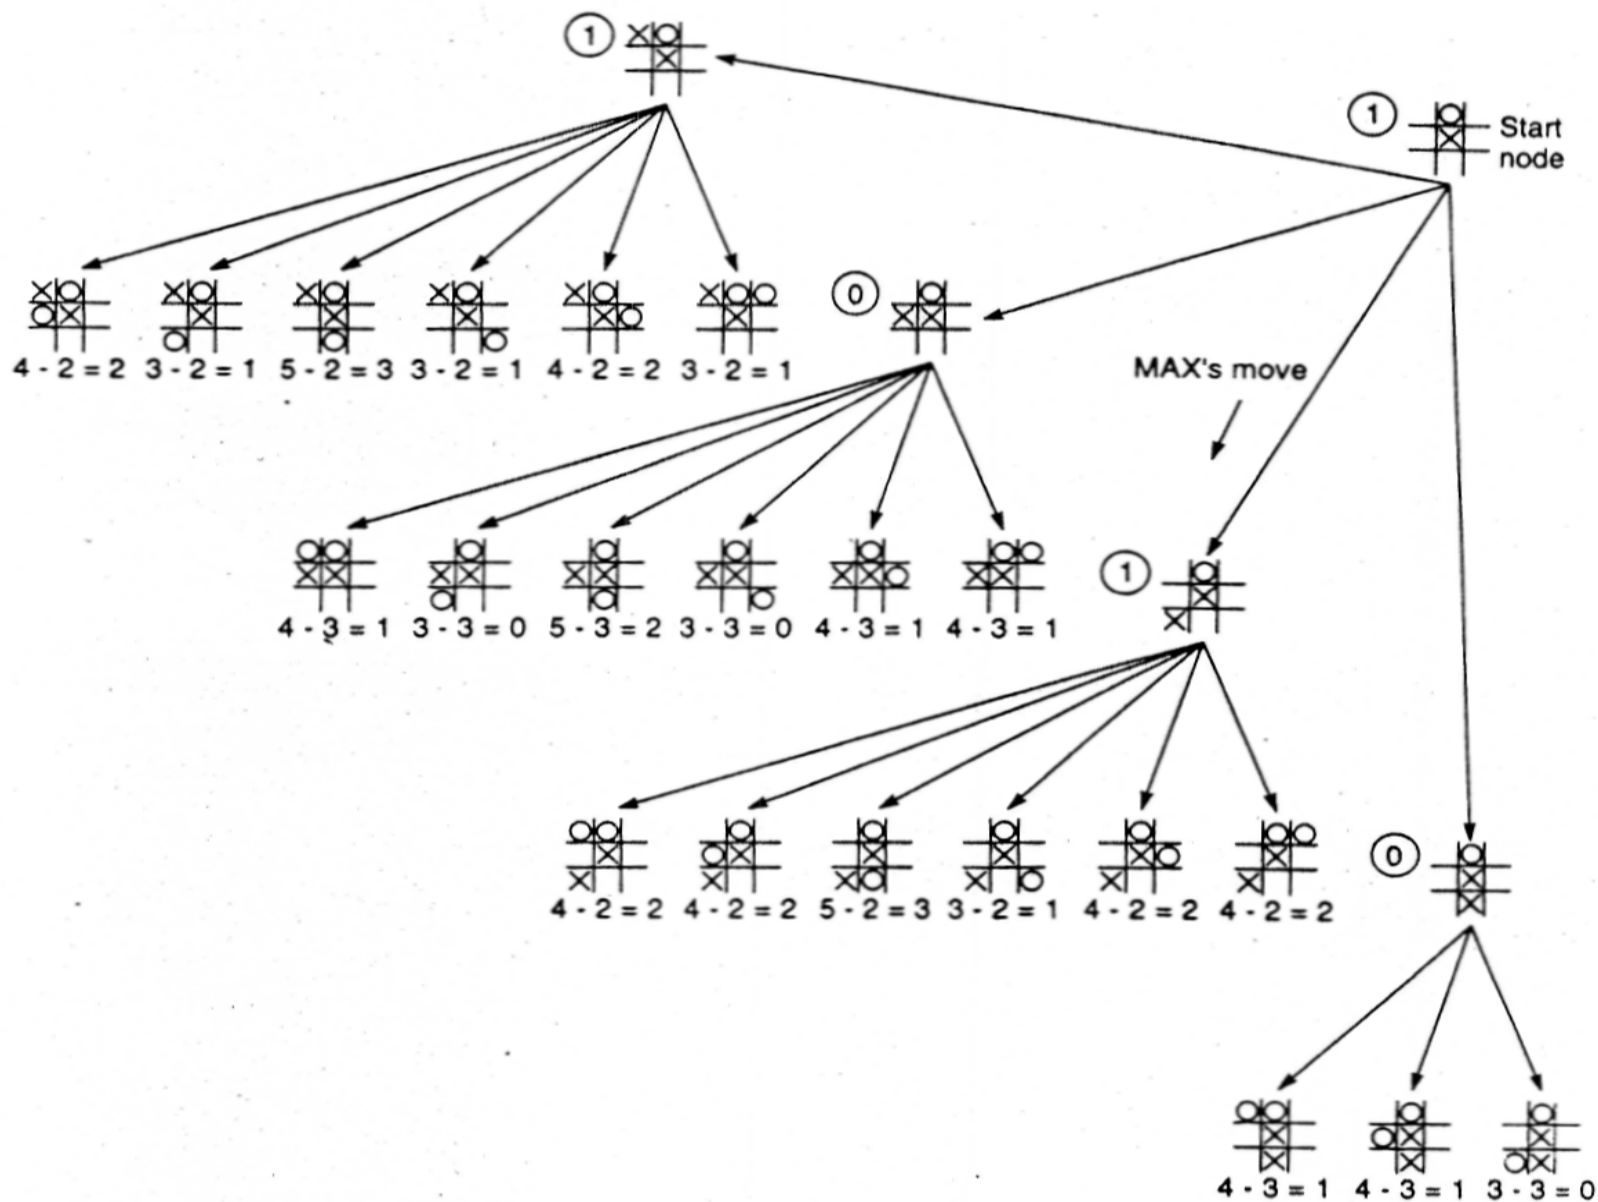
X	X	
O	O	O
	X	

$$e(p) = -\infty$$

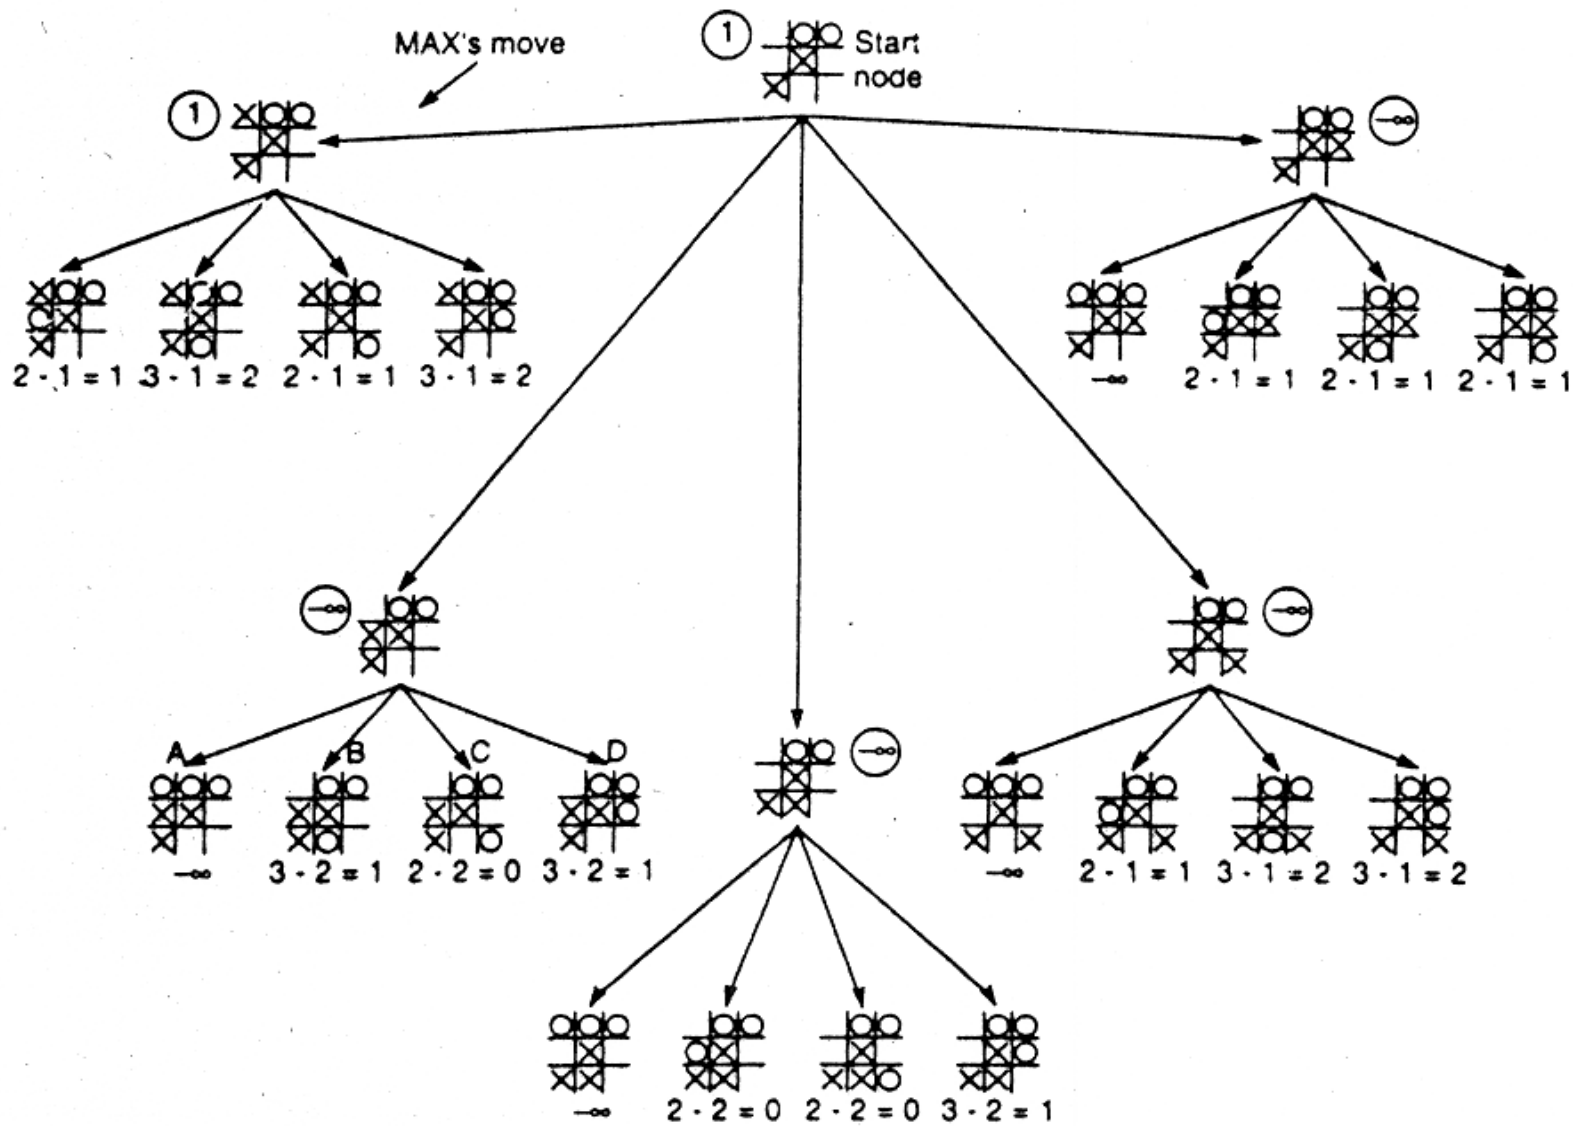
# Backup Values



**Figure 4.17** Two-ply minimax applied to the opening move of tic-tac-toe.



**Figure 4.18** Two-ply minimax applied to X's second move of tic-tac-toe.



**Figure 4.19** Two-ply minimax applied to X's move near end game.

# Alpha Beta Procedure

- Idea:
  - Do Depth first search to generate partial game tree.
  - Give static evaluation function to leaves.
  - compute bound on internal nodes.
- Alpha, Beta bounds:
  - Alpha value for Max node means that Max real value is at least alpha.
  - Beta for Min node means that Min can guarantee a value below Beta.
- Computation:
  - Alpha of a Max node is the maximum value of its seen children.
  - Beta of a Min node is the lowest value seen of its child node .

# When to Prune

## ❖ Pruning

- ❖ Below a Min node whose beta value is lower than or equal to the alpha value of a max ancestor.
- ❖ Below a Max node having an alpha value greater than or equal to the beta value of a min ancestor.

# The Alpha-Beta Procedure

- ❖ Now let us specify how to prune the Minimax tree in the case of a static evaluation function.
- ❖ Use two variables alpha (associated with MAX nodes) and beta (associated with MIN nodes).
- ❖ These variables contain the best (highest or lowest, resp.)  $e(p)$  value at a node  $p$  that has been found so far.
- ❖ Notice that alpha can never decrease, and beta can never increase.



# The Alpha-Beta Procedure

- ❖ There are two rules for terminating search:
- ❖ Search can be stopped below any MIN node having a beta value less than or equal to the alpha value of any of its MAX ancestors.
- ❖ Search can be stopped below any MAX node having an alpha value greater than or equal to the beta value of any of its MIN ancestors.

# alpha-beta pruning algorithm

## Procedure MinimaxAlphaBeta

**Input:** Position  $u$ , value  $\alpha$ , value  $\beta$


**Output:** Value at root

```
if (leaf( $u$ )) return  $Eval(p)$            ;; No successor, return evaluation
if (max-node( $u$ ))                        ;; MAX node
     $res \leftarrow \alpha$                  ;; Initialize result value
    for each  $v \in Succ(u)$                ;; Traverse successor list
         $val \leftarrow MinimaxAlphaBeta(v, res, \beta)$            ;; Recursion for  $\alpha$ 
         $res \leftarrow \max\{res, val\}$            ;; Take maximal value
        if ( $res \geq \beta$ )                 ;; Result exceeds threshold
            return  $res$                    ;; Propagate value
    else                                ;; MIN node
         $res \leftarrow \beta$                ;; Initialize result value
        for each  $v \in Succ(u)$            ;; Traverse successor list
             $val \leftarrow MinimaxAlphaBeta(v, \alpha, res)$        ;; Recursion for  $\beta$ 
             $res \leftarrow \min\{res, val\}$        ;; Take minimal value
            if ( $res \leq \alpha$ )           ;; Result exceeds threshold
                return  $res$                ;; Propagate value
return  $res$                              ;; Propagate value
```

## Algorithm 12.4


Minimax game tree search with  $\alpha\beta$ -pruning.

# $\alpha$ - $\beta$ pruning example



MAX


MIN




MAX knows that it can at least get "3" by playing this branch

MIN will choose "3", because it minimizes the utility (which is good for MIN)

# $\alpha$ - $\beta$ pruning example




## $\alpha$ - $\beta$ pruning example




MAX can get at least 14 in this branch  
so MAX will want to explore this branch more.

# $\alpha$ - $\beta$ pruning example



Should continue to explore

# $\alpha$ - $\beta$ pruning example



Search completed without exploring two branches


# Properties of $\alpha$ - $\beta$ pruning

- Pruning **does not** affect final result (it is exact).
- Good **move ordering** improves effectiveness of pruning (see last branch in example).
- The values at intermediate nodes may not be the same as the values computed by the minmax algorithm.




# The Alpha-Beta Procedure

Example:




# The Alpha-Beta Procedure

Example:




# The Alpha-Beta Procedure

Example:




# The Alpha-Beta Procedure

Example:




# The Alpha-Beta Procedure

Example:




Example:




# The Alpha-Beta Procedure

Example:



# The Alpha-Beta Procedure


Example:






# The Alpha-Beta Procedure

Example:




# The Alpha-Beta Procedure

Example:




## Example:




# The Alpha-Beta Procedure

Example:




## Example:




# The Alpha-Beta Procedure

## Example:




# The Alpha-Beta Procedure

Example:



# The Alpha-Beta Procedure


Example:






# The Alpha-Beta Procedure

Example:



# Nodes pruned by $\alpha$ - $\beta$ procedure

Another example



# The Alpha-Beta Procedure

- ❖ Can we estimate the **efficiency benefit** of the alpha-beta method?
- ❖ Suppose that there is a game that always allows a player to choose among **b** different moves, and we want to look **d** moves ahead.
- ❖ Then our search tree has  **$b^d$  leaves**.
- ❖ Therefore, if we do not use alpha-beta pruning, we would have to apply the static evaluation function  $N_d = b^d$  times.

# The Alpha-Beta Procedure

- ❖ Of course, the efficiency gain by the alpha-beta method always depends on the rules and the current configuration of the game.
- ❖ However, if we assume that new children of a node are explored in a particular order - those nodes  $p$  are explored first that will yield maximum values  $e(p)$  at depth  $d$  for MAX and minimum values for MIN - the number of nodes to be evaluated is:

# Timing Issues

- ❖ It is very difficult to predict for a given game situation how many operations a depth  $d$  look-ahead will require.
- ❖ Since we want the computer to respond within a certain amount of time, it is a good idea to apply the idea of **iterative deepening**.
- ❖ First, the computer finds the best move according to a **one-move look-ahead** search.
- ❖ Then, the computer determines the best move for a **two-move look-ahead**, and remembers it as the new best move.
- ❖ This is **continued** until the time runs out. Then the currently remembered best move is executed.

# How to Find Static Evaluation Functions

- ❖ Often, a static evaluation function  $e(p)$  first computes an appropriate **feature vector  $f(p)$**  that contains information about features of the current game configuration that are important for its evaluation.
- ❖ There is also a **weight vector  $w(p)$**  that indicates the weight (= importance) of each feature for the assessment of the current situation.
- ❖ Then  $e(p)$  is simply computed as the **scalar product** of  $f(p)$  and  $w(p)$ .
- ❖ Both the identification of the most relevant features and the correct estimation of their relative importance are **crucial** for the strength of a game-playing program.

## Evaluation function for chess

- For games like chess, typically linear weighted sum of features
- $$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.  $w_1 = 9$  with
- $f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$
- Weights 9 for queen, 5 for rook, 3 for bishop and knight and 1 for pawn – suggested by Shannon is still widely used.

## How to Find Static Evaluation Functions

- ❖ Once we have found suitable features, the weights can be **adapted algorithmically**.
- ❖ This can be achieved, for example, with a **neural network**.
- ❖ So the challenge is in extracting the most informative features from a game configuration.








# Forward Pruning & Lookup tables

- Humans don't consider all possible moves.
- Can we prune certain branches immediately?
- Use estimates (from past experience) of the uncertainty in the estimate of the node's value and uses that to decide if a node can be pruned.
- Instead of search one can also store game states.
- Openings in chess are played from a library
- Endgames have often been solved and stored as well.

# Chance Games


## Backgammon



element of chance

# Expected Minimax (Expectimax)

Again, the tree is constructed bottom-up.



Now we have even more nodes to search!

