

Alpha Beta Procedure

- Idea:
 - Do Depth first search to generate partial game tree.
 - Give static evaluation function to leaves.
 - compute bound on internal nodes.
- Alpha, Beta bounds:
 - Alpha value for Max node means that Max real value is at least alpha.
 - Beta for Min node means that Min can guarantee a value below Beta.
- Computation:
 - Alpha of a Max node is the maximum value of its seen children.
 - Beta of a Min node is the lowest value seen of its child node .

When to Prune

❖ Pruning

- ❖ Below a Min node whose beta value is lower than or equal to the alpha value of its ancestors.
- ❖ Below a Max node having an alpha value greater than or equal to the beta value of any of its Min nodes ancestors.

The Alpha-Beta Procedure

The Alpha-Beta Procedure

- ❖ Now let us specify how to prune the Minimax tree in the case of a static evaluation function.

The Alpha-Beta Procedure

- ❖ Now let us specify how to prune the Minimax tree in the case of a static evaluation function.
- ❖ Use two variables α (associated with MAX nodes) and β (associated with MIN nodes).

The Alpha-Beta Procedure

- ❖ Now let us specify how to prune the Minimax tree in the case of a static evaluation function.
- ❖ Use two variables alpha (associated with MAX nodes) and beta (associated with MIN nodes).
- ❖ These variables contain the best (highest or lowest, resp.) $e(p)$ value at a node p that has been found so far.

The Alpha-Beta Procedure

- ❖ Now let us specify how to prune the Minimax tree in the case of a static evaluation function.
- ❖ Use two variables alpha (associated with MAX nodes) and beta (associated with MIN nodes).
- ❖ These variables contain the best (highest or lowest, resp.) $e(p)$ value at a node p that has been found so far.
- ❖ Notice that alpha can never decrease, and beta can never increase.

The Alpha-Beta Procedure

The Alpha-Beta Procedure

- ❖ There are two rules for terminating search:

The Alpha-Beta Procedure

- ❖ There are two rules for terminating search:
- ❖ Search can be stopped below any MIN node having a beta value less than or equal to the alpha value of any of its MAX ancestors.

The Alpha-Beta Procedure

- ❖ There are two rules for terminating search:
- ❖ Search can be stopped below any MIN node having a beta value less than or equal to the alpha value of any of its MAX ancestors.
- ❖ Search can be stopped below any MAX node having an alpha value greater than or equal to the beta value of any of its MIN ancestors.

The Alpha-Beta Procedure

- ❖ There are two rules for terminating search:
- ❖ Search can be stopped below any MIN node having a beta value less than or equal to the alpha value of any of its MAX ancestors.
- ❖ Search can be stopped below any MAX node having an alpha value greater than or equal to the beta value of any of its MIN ancestors.

Alpha-Beta procedure

```
def value(state,  $\alpha$ ,  $\beta$ ):  
    if (state is a max node):  
        max-value((state,  $\alpha$ ,  $\beta$ )  
    else:  
        min-value((state,  $\alpha$ ,  $\beta$ )
```


α : MAX's best option on path to root

β : MIN's best option on path to root


```
def max-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = -\infty$   
    if leaf(state):  
        return value(state)  
    for each successor of state:  
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \geq \beta$  return  $v$   
         $\alpha = \max(\alpha, v)$   
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):  
    initialize  $v = +\infty$   
    if leaf(state):  
        return value(state)  
    for each successor of state:  
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$   
        if  $v \leq \alpha$  return  $v$   
         $\beta = \min(\beta, v)$   
    return  $v$ 
```


Pruning the tree Example 1



Pruning the tree Example




α - β pruning example 2



MAX


MIN




MAX knows that it can at least get "3" by playing this branch

MIN will choose "3", because it minimizes the utility (which is good for MIN)

α - β pruning example 2




α - β pruning example 2




MAX can get at least 14 in this branch
so MAX will want to explore this branch more.

α - β pruning example 2



Should continue to explore


α - β pruning example 2



Search completed without exploring two branches


The Alpha-Beta Procedure Example 3

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:



α - β pruning quiz



Properties of α - β pruning

- Pruning **does not** affect final result (it is exact).
- Good **move ordering** improves effectiveness of pruning (see last branch in example).
- The values at intermediate nodes may not be the same as the values computed by the minmax algorithm.

The Alpha-Beta Procedure

The Alpha-Beta Procedure

- ❖ What is the **efficiency benefit** of the alpha-beta method? (It is no better than min-max procedure of the same depth in terms of the quality of the move.)

The Alpha-Beta Procedure

- ❖ What is the **efficiency benefit** of the alpha-beta method? (It is no better than min-max procedure of the same depth in terms of the quality of the move.)
- ❖ Suppose that there is a game that always allows a player to choose among **b** different moves, and we want to look **d** moves ahead.

The Alpha-Beta Procedure

- ❖ What is the **efficiency benefit** of the alpha-beta method? (It is no better than min-max procedure of the same depth in terms of the quality of the move.)
- ❖ Suppose that there is a game that always allows a player to choose among **b** different moves, and we want to look **d** moves ahead.
- ❖ Then our search tree has **b^d leaves**.

The Alpha-Beta Procedure

- ❖ What is the **efficiency benefit** of the alpha-beta method? (It is no better than min-max procedure of the same depth in terms of the quality of the move.)
- ❖ Suppose that there is a game that always allows a player to choose among **b** different moves, and we want to look **d** moves ahead.
- ❖ Then our search tree has **b^d leaves**.
- ❖ If we do not use alpha-beta pruning, we would have to apply the static evaluation function $N_d = b^d$ times.

The Alpha-Beta Procedure

The Alpha-Beta Procedure

- ❖ if we assume that new children of a node are explored in the “most beneficial” order - those nodes p are explored first that will yield maximum values $e(p)$ at depth d for MAX and minimum values for MIN - the number of nodes to be evaluated is:

Timing Issues

- ❖ It is very difficult to predict for a given game situation how many operations a depth d look-ahead will require.

Timing Issues

- ❖ It is very difficult to predict for a given game situation how many operations a depth d look-ahead will require.
- ❖ Since we want the computer to respond within a certain amount of time, it is a good idea to apply the idea of **iterative deepening**.

Timing Issues

- ❖ It is very difficult to predict for a given game situation how many operations a depth d look-ahead will require.
- ❖ Since we want the computer to respond within a certain amount of time, it is a good idea to apply the idea of **iterative deepening**.
- ❖ First, the computer finds the best move according to a **one-move look-ahead** search.

Timing Issues

- ❖ It is very difficult to predict for a given game situation how many operations a depth d look-ahead will require.
- ❖ Since we want the computer to respond within a certain amount of time, it is a good idea to apply the idea of **iterative deepening**.
- ❖ First, the computer finds the best move according to a **one-move look-ahead** search.
- ❖ Then, the computer determines the best move for a **two-move look-ahead**, and remembers it as the new best move.

Timing Issues

- ❖ It is very difficult to predict for a given game situation how many operations a depth d look-ahead will require.
- ❖ Since we want the computer to respond within a certain amount of time, it is a good idea to apply the idea of **iterative deepening**.
- ❖ First, the computer finds the best move according to a **one-move look-ahead** search.
- ❖ Then, the computer determines the best move for a **two-move look-ahead**, and remembers it as the new best move.
- ❖ This is **continued** until the time runs out. Then the currently remembered best move is executed.

How to Find Static Evaluation Functions

- ❖ Often, a static evaluation function $e(p)$ first computes an appropriate **feature vector $f(p)$** that contains information about features of the current game configuration that are important for its evaluation.

How to Find Static Evaluation Functions

- ❖ Often, a static evaluation function $e(p)$ first computes an appropriate **feature vector $f(p)$** that contains information about features of the current game configuration that are important for its evaluation.
- ❖ There is also a **weight vector $w(p)$** that indicates the weight (= importance) of each feature for the assessment of the current situation.

How to Find Static Evaluation Functions

- ❖ Often, a static evaluation function $e(p)$ first computes an appropriate **feature vector $f(p)$** that contains information about features of the current game configuration that are important for its evaluation.
- ❖ There is also a **weight vector $w(p)$** that indicates the weight (= importance) of each feature for the assessment of the current situation.
- ❖ Then $e(p)$ is simply computed as the **scalar product** of $f(p)$ and $w(p)$.

How to Find Static Evaluation Functions

- ❖ Often, a static evaluation function $e(p)$ first computes an appropriate **feature vector $f(p)$** that contains information about features of the current game configuration that are important for its evaluation.
- ❖ There is also a **weight vector $w(p)$** that indicates the weight (= importance) of each feature for the assessment of the current situation.
- ❖ Then $e(p)$ is simply computed as the **scalar product** of $f(p)$ and $w(p)$.
- ❖ Both the identification of the most relevant features and the correct estimation of their relative importance are **crucial** for the strength of a game-playing program.

Evaluation function for chess

- For games like chess, typically linear weighted sum of features
- $$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g. $w_1 = 9$ with
- $f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$
- Weights 9 for queen, 5 for rook, 3 for bishop and knight and 1 for pawn – suggested by Shannon is still widely used.

How to Find Static Evaluation Functions

- ❖ Once we have found suitable features, the weights can be **adapted algorithmically**.

How to Find Static Evaluation Functions

- ❖ Once we have found suitable features, the weights can be **adapted algorithmically**.
- ❖ This can be achieved, for example, with a **neural network**.

How to Find Static Evaluation Functions

- ❖ Once we have found suitable features, the weights can be **adapted algorithmically**.
- ❖ This can be achieved, for example, with a **neural network**.
- ❖ So the challenge is in extracting the most informative features from a game configuration.

How to Find Static Evaluation Functions


- ❖ Once we have found suitable features, the weights can be **adapted algorithmically**.
- ❖ This can be achieved, for example, with a **neural network**.
- ❖ So the challenge is in extracting the most informative features from a game configuration.

Forward Pruning & Lookup tables

- Humans don't consider all possible moves.
- Can we prune certain branches immediately?
- Use estimates (from past experience) of the uncertainty in the estimate of the node's value and uses that to decide if a node can be pruned.
- Instead of search one can also store game states.
- Openings in chess are played from a library
- Endgames have often been solved and stored as well.

Chance Games


Backgammon



element of chance

Expected Minimax (Expectimax)


Again, the tree is constructed bottom-up.



Now we have even more nodes to search!

Multi-Agent Utilities

- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
 - Terminals have utility tuples
 - Node values are also utility tuples
 - Each player maximizes its own component
 - Can give rise to cooperation and competition dynamically...



Multi-Agent Utilities

- What if the game is not zero-sum, or has multiple players?
- Generalization of minimax:
 - Terminals have utility tuples
 - Node values are also utility tuples
 - Each player maximizes its own component
 - Can give rise to cooperation and competition dynamically...

