

# Reinforcement Learning

## Markov Decision Processes



[These slides adapted from Dan Klein and Pieter Abbeel]


# Non-Deterministic Search

---

- Recall search problems in which the goal was to find the least cost path from the starting state to a goal state.
- In a stochastic (or non-deterministic) search problem, there is more than one possible next state when a particular action is taken.
- $T(s, a, s')$  is the probability that the state  $s'$  is reached when action  $a$  is taken in state  $s$ .


# Example: Grid World

- A maze-like problem
  - The agent lives in a grid
  - Walls block the agent's path
- Noisy movement: actions do not always go as planned
  - 80% of the time, the action North takes the agent North (if there is no wall there) If there is a wall in the direction the agent would have been taken, the agent stays put with prob = 0.8
  - 10% of the time, North takes the agent West; 10% East
- The agent receives rewards each time step
  - Small “living” reward each step (usually negative)
  - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards




# Grid World Actions

Deterministic Grid World




Stochastic Grid World



# Markov Decision Processes

- An MDP is defined by:
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Probability that  $a$  from  $s$  leads to  $s'$ , i.e.,  $P(s' | s, a)$
    - Also called the model or the dynamics
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s')$
  - A **start state**
  - One or more **terminal states**



# The significance of the name MDP

---

- “Markov” generally means that given the present state, the future and the past are independent
- For Markov decision processes, “Markov” means action outcomes depend only on the current state

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$$

=

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$


- This is just like search, where the successor function could only depend on the current state (not the history)



Andrey Markov  
(1856-1922)


# Policies

- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal **policy  $\pi^*: S \rightarrow A$** 
  - A policy  $\pi$  gives an action for each state
  - An optimal policy is one that maximizes expected utility if followed
  - An explicit policy defines a reflex agent
- policy is what we needed in a two-player game as well (since a fixed path can't be predicted in advance.)




Optimal policy when  $R(s, a, s') = -0.03$   
for all non-terminals  $s$


# Optimal Policies




$R(s) = -0.01$



$R(s) = -0.03$



$R(s) = -0.4$




$R(s) = -2.0$



# Utilities of Sequences


---



# Utilities of Sequences

---

- What preferences should an agent have over reward sequences?
- More or less?  $[1, 2, 2]$  or  $[2, 3, 4]$
- Now or later?  $[0, 0, 1]$  or  $[1, 0, 0]$



# Discounting

---

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution: values of rewards decay exponentially



1

Worth Now



$\gamma$

Worth Next Step



$\gamma^2$

Worth In Two Steps

# Discounting

- How to discount?


- Each time we descend a level, we multiply in the discount once

- Why discount?

- Think of it as a gamma chance of ending the process at every step
- Also helps our algorithms converge

- Example: discount of 0.5

- $U([1,2,3]) = 1*1 + 0.5*2 + 0.25*3$
- $U([1,2,3]) < U([3,2,1])$



# Quiz: Discounting

- Given:

10				1
a	b	c	d	e

- Actions: East, West, and Exit (only available in exit states a, e)
- Transitions: deterministic

- Quiz 1: For  $\gamma = 1$ , what is the optimal policy?

10	<-	<-	<-	1
----	----	----	----	---

- Quiz 2: For  $\gamma = 0.1$ , what is the optimal policy?

10	<-	<-	->	1
----	----	----	----	---

- Quiz 3: For which  $\gamma$  are West and East equally good when in state d?

$$10\gamma = \gamma^3$$

# Infinite Utilities?!

---

- Problem: What if the game lasts forever? Do we get infinite rewards?

- Solutions:

- Finite horizon: (similar to depth-limited search)
  - Terminate episodes after a fixed T steps (e.g. life)
  - Gives nonstationary policies ( $\pi$  depends on time left)


- Discounting: use  $0 < \gamma < 1$

- $$\text{Sum} U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max}/(1 - \gamma)$$

- Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like “overheated” for racing)


# Example: Racing

---




# Example: Racing

- A robot car wants to travel far, quickly
- Three states: **Cool**, **Warm**, Overheated
- Two actions: *Slow*, *Fast*
- Going faster gets double reward






# Racing Search Tree



# MDP Search Trees


- Each MDP state projects an expectimax-like search tree



# Recap: Defining MDPs


---


- Markov decision processes:
  - Set of states  $S$
  - Start state  $s_0$
  - Set of actions  $A$
  - Transitions  $P(s' | s, a)$  (or  $T(s, a, s')$ )
  - Rewards  $R(s, a, s')$  (and discount  $\gamma$ )
- MDP quantities so far:
  - Policy = Choice of action for each state
  - Utility = sum of (discounted) rewards




# Solving MDPs

---






# Racing Search Tree




# Racing Search Tree

- We're doing way too much work with expectimax!
- Problem: States are repeated
  - Idea: Only compute needed quantities once
- Problem: Tree goes on forever
  - Idea: Do a depth-limited computation, but with increasing depths until change is small
  - Note: deep parts of the tree eventually don't matter if  $\gamma < 1$



# Optimal Quantities

- The value (utility) of a state  $s$ :  
 $V^*(s)$  = expected utility starting in  $s$  and acting optimally
- The value (utility) of a q-state  $(s,a)$ :  
 $Q^*(s,a)$  = expected utility starting out having taken action  $a$  from state  $s$  and (thereafter) acting optimally
- The optimal policy:  
 $\pi^*(s)$  = optimal action from state  $s$






# Gridworld v Values after 100 iterations



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# Snapshot of Demo – Gridworld Q Values



Noise = 0.2  
Discount = 0.9  
Living reward = 0


# Values of States

---

- Recursive definition of value:

$$V^*(s) = \max_a Q^*(s, a)$$





$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



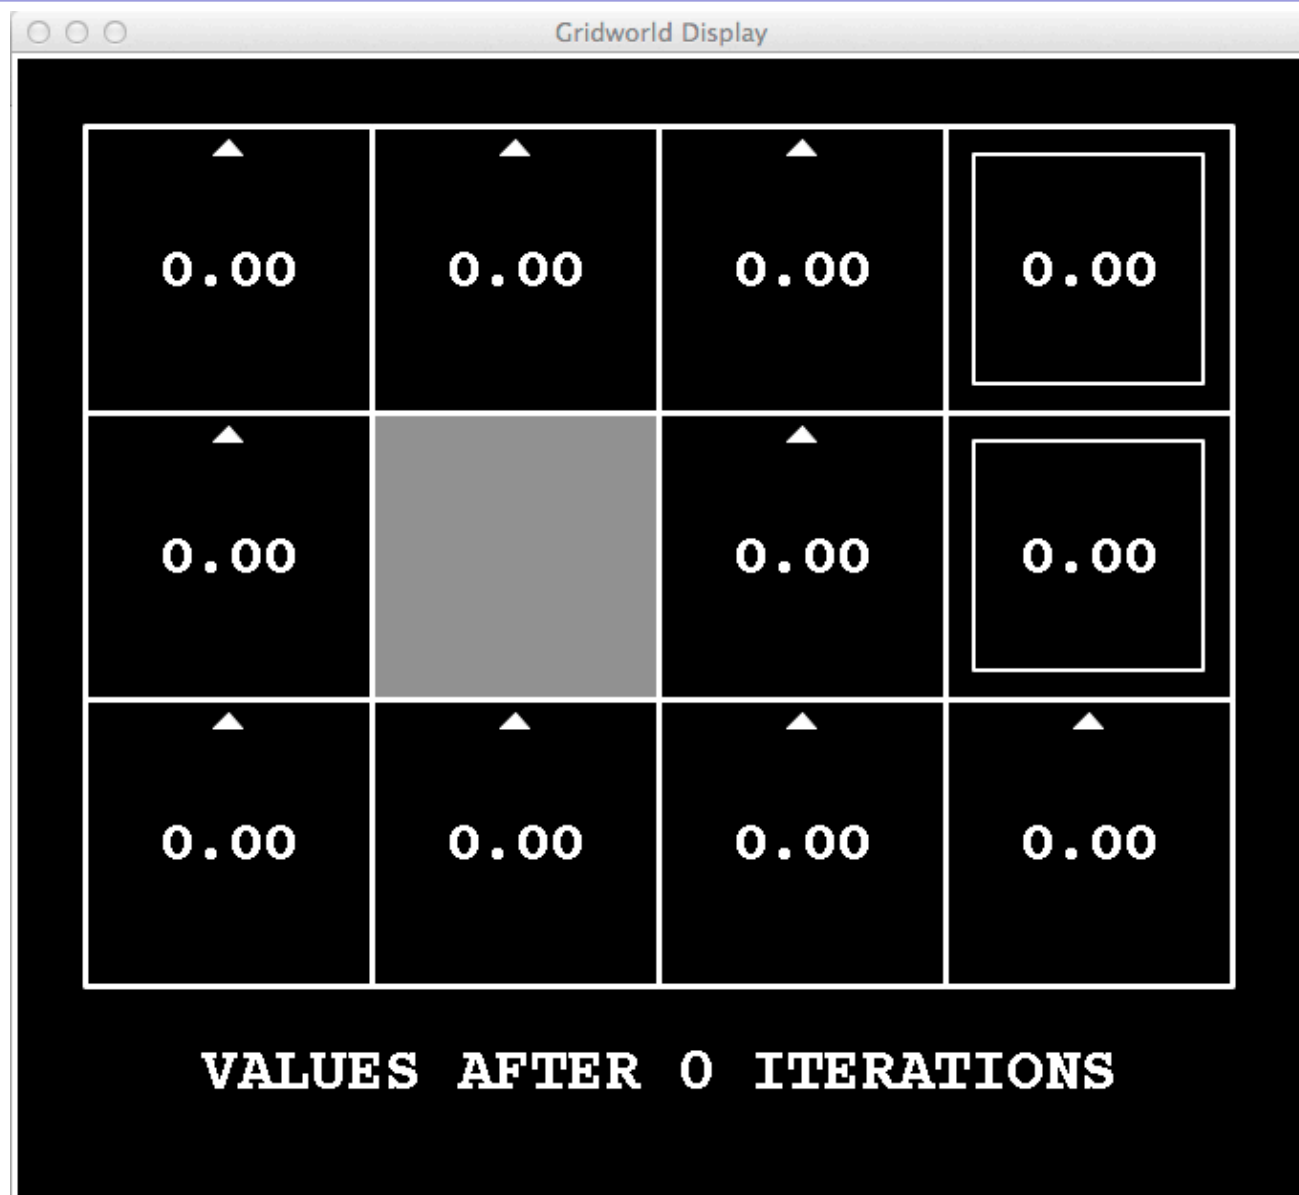
$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

# Time-Limited Values

- Key idea: time-limited values
- Define  $V_k(s)$  to be the optimal value of  $s$  if the game ends in  $k$  more time steps
  - Equivalently, it's what a depth- $k$  expectimax would give from  $s$

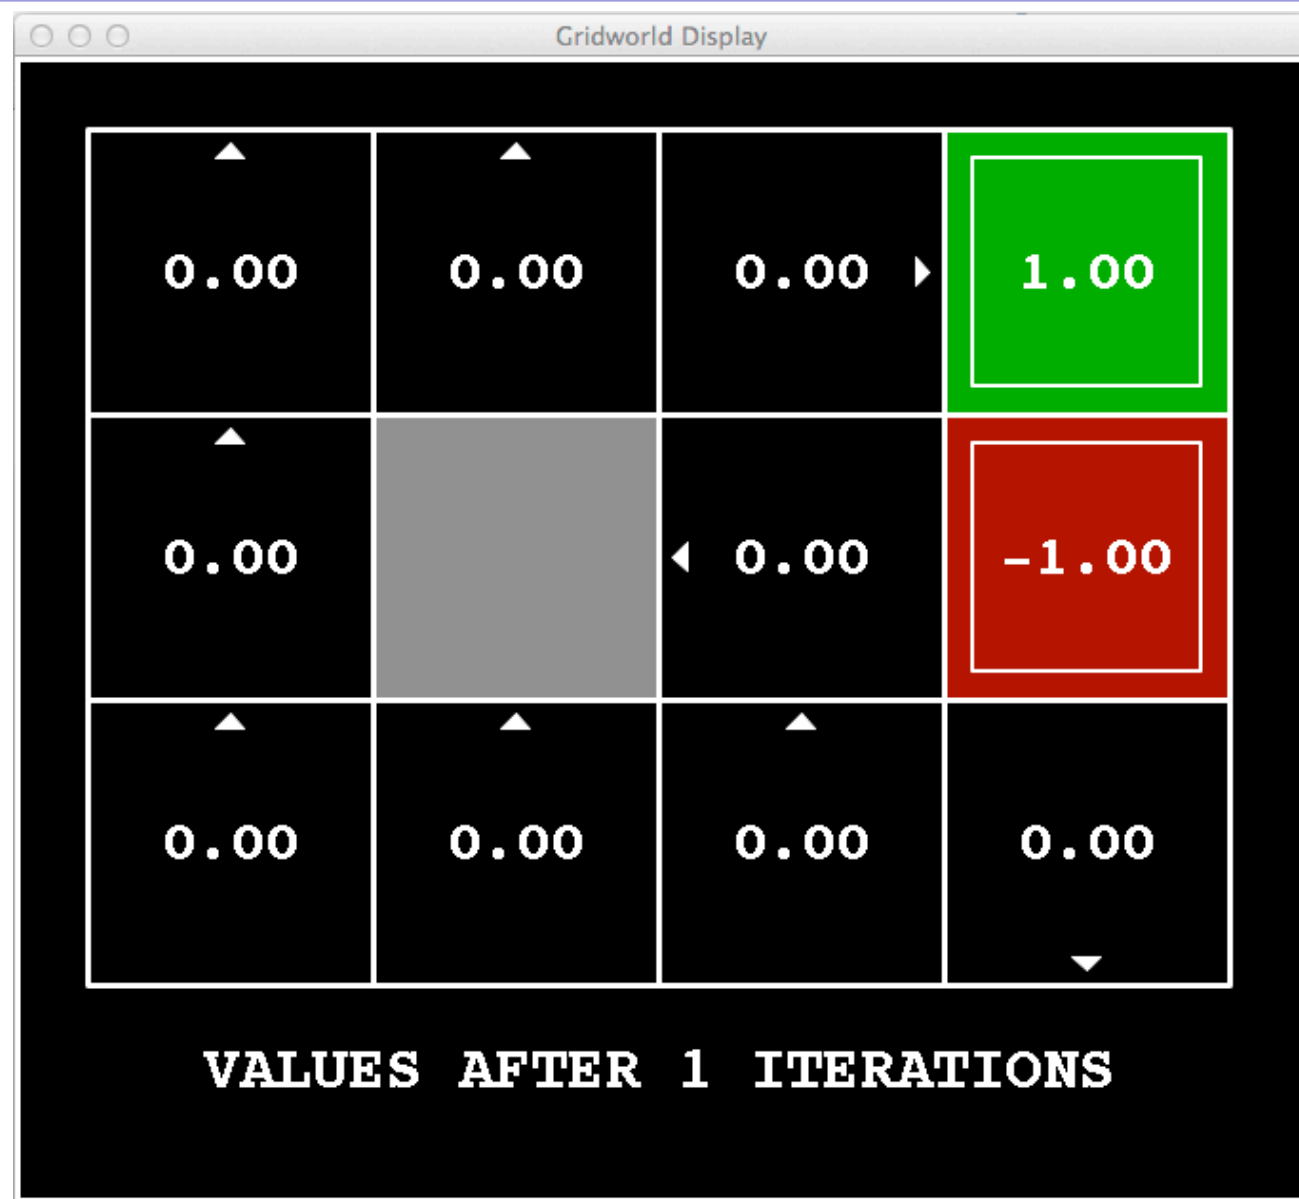


$k=0$



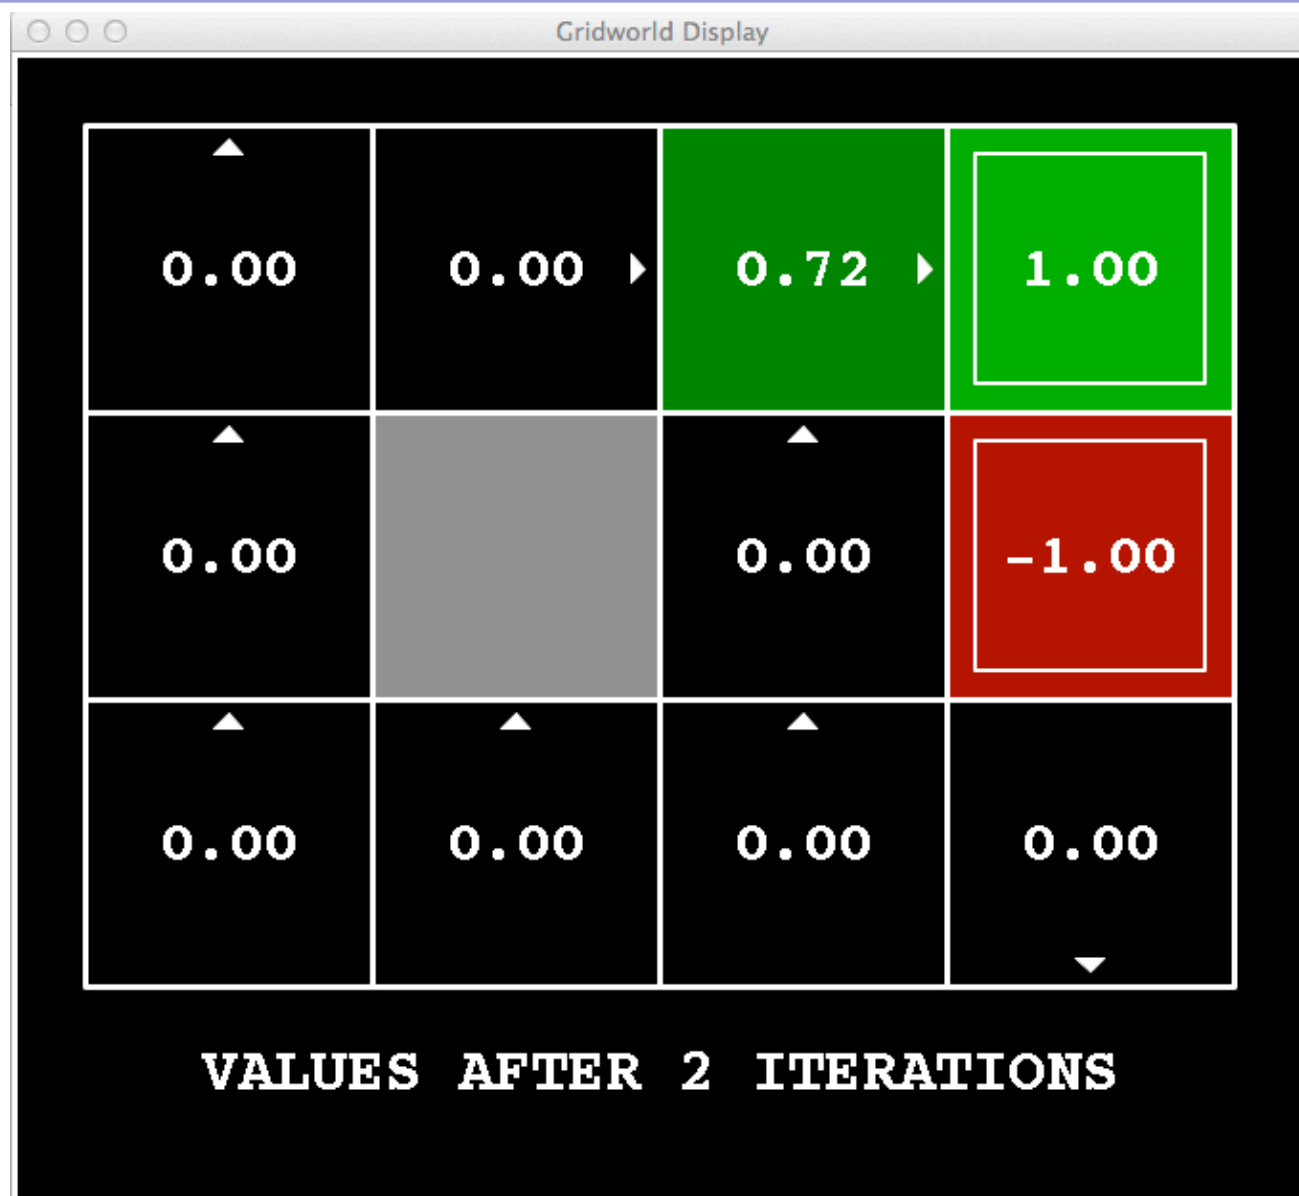
Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=1$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=2$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=3$



Noise = 0.2  
Discount = 0.9  
Living reward = 0



$k=4$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=5$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=6$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=7$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=8$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

$k=9$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# $k=10$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# $k=11$



Noise = 0.2  
Discount = 0.9  
Living reward = 0



# k=12



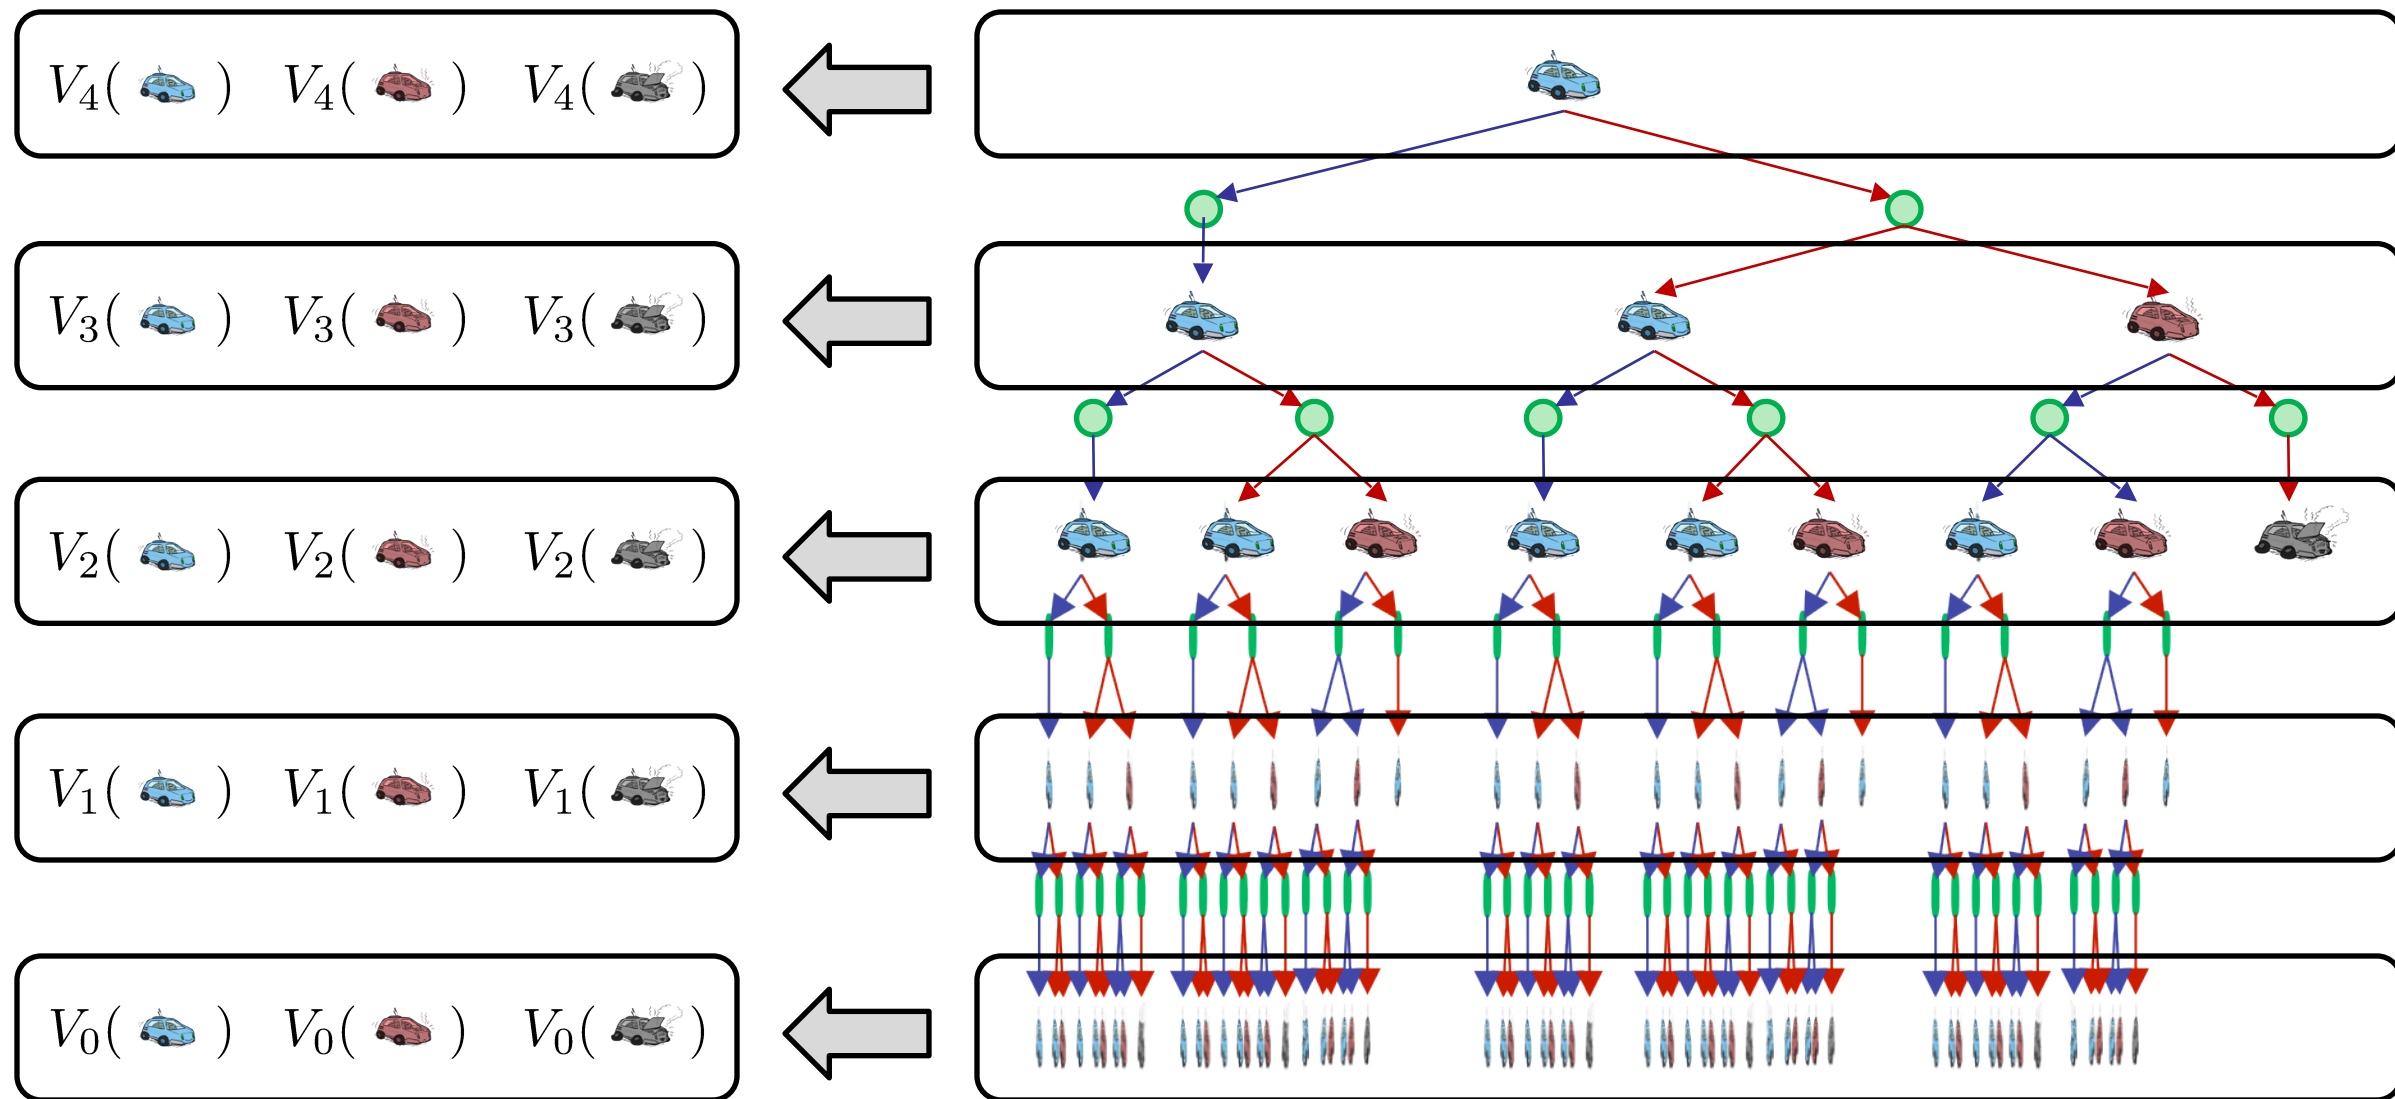
Noise = 0.2  
Discount = 0.9  
Living reward = 0

# $k=100$



Noise = 0.2  
Discount = 0.9  
Living reward = 0

# Computing Time-Limited Values

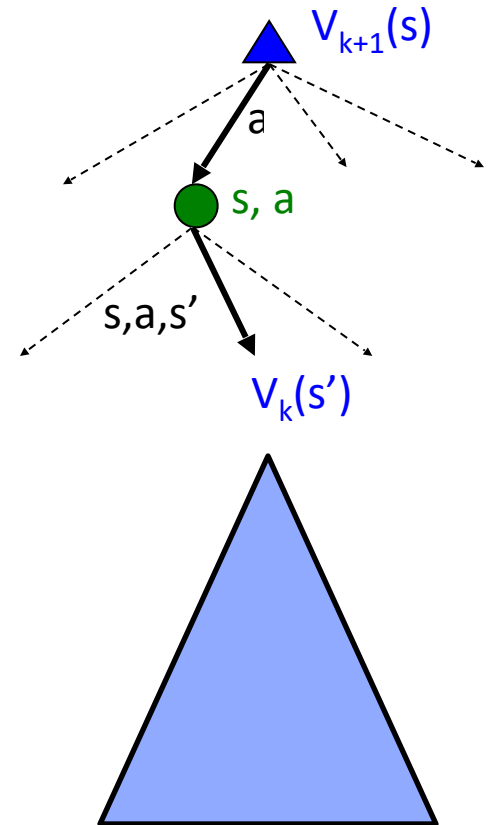


# Value Iteration

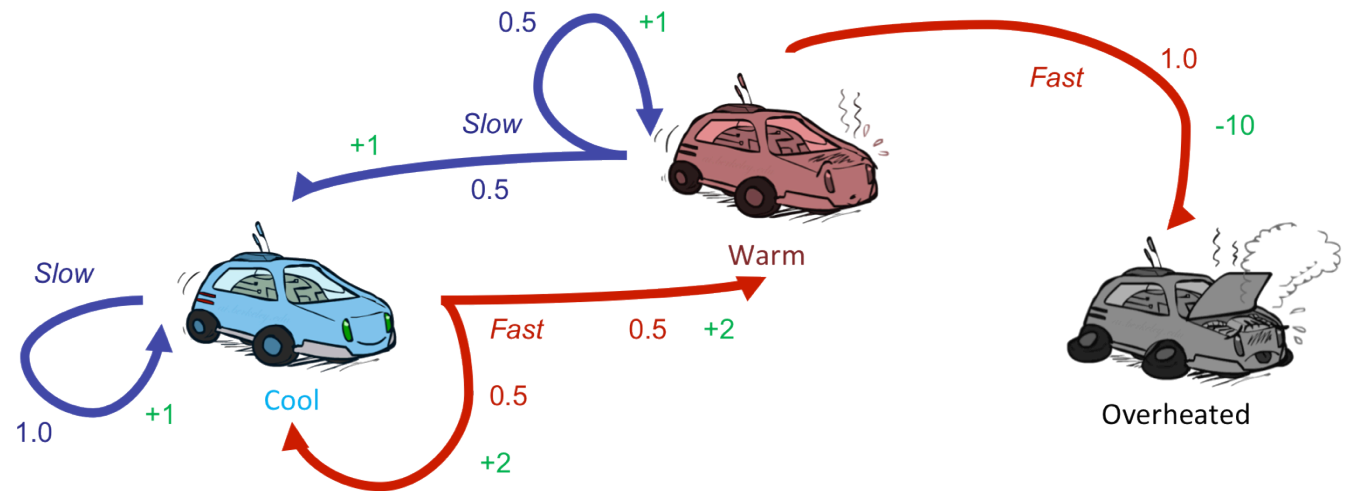

- Start with  $V_0(s) = 0$ : no time steps left means an expected reward sum of zero
- Given vector of  $V_k(s)$  values, do one ply of expectimax from each state:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Repeat until convergence
- Complexity of each iteration:  $O(S^2A)$
- Theorem: will converge to unique optimal values
  - Basic idea: approximations get refined towards optimal values
  - Policy may converge long before values do






# Example: Value Iteration




Assume no discount!

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

# Example: Value Iteration




			
$V_2$			
$V_1$	2	S: $.5*1+.5*1=1$ F: -10	
$V_0$	0	0	0




*Assume no discount!*

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

# Example: Value Iteration

			
$V_2$			
$V_1$	2	1	0
$V_0$	0	0	0



Assume no discount!

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

# Example: Value Iteration



$V_2$

S:  $1+2=3$

F:  $.5*(2+2)+.5*(2+1)=3.5$

$V_1$

2

1


0

$V_0$

0

0

0







*Assume no discount!*

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$



# Example: Value Iteration

			
$V_2$	3.5	2.5	0
$V_1$	2	1	0
$V_0$	0	0	0



Assume no discount!

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

# Convergence\*

- How do we know the  $V_k$  vectors are going to converge?
- Case 1: If the tree has maximum depth  $M$ , then  $V_M$  holds the actual untruncated values
- Case 2: If the discount is less than 1
  - Sketch: For any state  $V_k$  and  $V_{k+1}$  can be viewed as depth  $k+1$  expectimax results in nearly identical search trees
  - The difference is that on the bottom layer,  $V_{k+1}$  has actual rewards while  $V_k$  has zeros
  - That last layer is at best all  $R_{\text{MAX}}$
  - It is at worst  $R_{\text{MIN}}$
  - But everything is discounted by  $\gamma^k$  that far out
  - So  $V_k$  and  $V_{k+1}$  are at most  $\gamma^k \max |R|$  different
  - So as  $k$  increases, the values converge

